PV and Battery Trading in MGP and MI Auctions
Current Mathematical Model and Execution Workflow

`bess_pv_trading` repository

Implementation reviewed on 22 October 2026

Abstract

This document describes the implemented optimization and execution system for a single photovoltaic farm and an AC-coupled battery sharing one point of interconnection. It distinguishes auction plans, immutable market commitments, physical feedback execution, and ex-post settlement. The shared Gurobi formulation supports a risk-neutral scenario tree. The current historical January experiment instead uses one deterministic rolling weather forecast and perfect auction-price information, followed by execution against measured AC production. A separate continuous experiment applies daily planning chronologically while carrying actual battery state, outstanding orders and operator fee history across days. Separate full-information PV+BESS and PV-only optimizations quantify an ex-post operating value of storage. Equations below follow the current source code; they are not a proposal or a claim of a fully causal historical trading strategy.

Contents

1 Scope and implementation map **3**

2 Time, units, signs and information **4**

2.1 Delivery intervals and auction products 4

2.2 Notation and electrical basis 4

2.3 Forecast vintages and observability 5

3 PV power from weather forecasts **6**

3.1 Irradiance and orientation 6

3.2 Temperature, DC output and AC conversion 6

3.3 Chronological physical PV calibration 7

3.3.1 How the installed 2025 parameter sets were selected 8

3.3.2 Average, minimum and maximum errors against the previous model 9

4 Optimization at an auction decision **10**

4.1 Scenario tree and nonanticipativity 10

4.2 Physical variables and balances 10

4.3 Auction quantities and physical delivery 11

4.4 Current objective and sunk cash flows . . . 12
4.5 Market transaction fees and operator monthly volume . . . 12
4.6 What is free at each of the four decisions . . . 13

5 Battery degradation models **14**
5.1 Current AC throughput approximation . . . 14
5.2 Optional segmented depth-dependent aging . . . 14

6 Submission, publication and reoptimization workflow **15**

7 Physical feedback with realized PV **15**
7.1 Feasible interval power range . . . 15
7.2 Grid-tracking policy . . . 16
7.3 PV source allocation for reporting . . . 17

8 Settlement and expected-versus-realized accounting **17**

9 Value of reoptimization through frozen-plan replay **18**

10 Perfect-information schedules and the PV-only baseline **19**
10.1 PV+BESS versus the feedback controller . . . 19
10.2 PV-only asset counterfactual . . . 19

11 Configurable participation and matched-day profile comparison **20**
11.1 Sequential current-auction-only baseline . . . 21
11.2 Operational fallback after an infeasible auction decision . . . 21

12 Continuous operation with daily optimization horizons **23**
12.1 Delivery books and the chronological event sequence . . . 23
12.2 Forecast-based projection of the next initial state . . . 23
12.3 Daily free decisions, initial condition and final target . . . 24
12.4 Independent trajectories, accounting and restart . . . 24

13 Price Forecasting Status **25**

14 Optional optimized-imbalance formulation **26**

15 Data provenance, outputs and reproducibility **26**

16 Current limitations and extension boundaries **27**

1 Scope and implementation map

One farm has PV generation, an AC-coupled BESS, and a common AC grid connection. There is no local electricity demand. PV can be exported, stored, or curtailed; the battery can charge from PV or grid imports and discharge toward the grid. The market sequence is MGP (called DA in code), MI-A1, MI-A2, and MI-A3 (IDA1, IDA2, IDA3). Every submitted signed quantity is assumed to execute in full at the auction clearing price. Bid prices, acceptance uncertainty, market impact, and continuous intraday trading are outside this formulation.

The implementation has three distinct objects:

1. An optimization plan, which includes the current actionable auction order, future auction recourse, and a battery/PV schedule over the remaining horizon.
2. An immutable commitment ledger, updated at submission and publication.
3. A physical trajectory produced by a feedback/protection policy using realized interval-average PV. Its energy, wear, and deviations are settled separately from the optimization objective.

Source	Responsibility
`src/domain.py`	Validated grids, assets, scenarios, information nodes, states and solver contracts.
`src/optimization/model.py`	Physical, trading, product-resolution, optional imbalance constraints and objective.
`src/optimization/degradation.py`	Throughput coefficient and segmented cycle-depth model.
`src/optimization/validation.py`	Numerical validation independent of the solver constraint objects.
`src/transaction_fees.py`	Pure-Python fee integration, dated tariffs and operator monthly state.
`src/optimization/transaction_fees.py`	Flat-band LP costs and exact progressive ABS/PWL constraints.
`src/simulation/synthetic.py`	Shared event replay and auction submission/publication accounting.
`src/simulation/replay.py`	Residual scenario conditioning and bounded physical execution.
`src/simulation/gme_oracle.py`	Manual-data historical price oracle, rolling PV and full-information asset comparisons.
`src/simulation/continuous.py`	Consecutive delivery-day books, shadow boundary projections and uninterrupted execution.
`src/reporting/continuous_outputs.py`	Continuous trajectory accounting, paired comparisons and native solve timing.
`src/simulation/settlement.py`	Exact-grid actual settlement and expected/realized KPI reconciliation.
`src/forecasting/pv.py`	Weather-to-AC PV conversion.
`src/markets.py`	Configurable local auction gates, publication events and DST-aware delivery grids.

2 Time, units, signs and information

2.1 Delivery intervals and auction products

A local delivery day is partitioned into physical intervals $\mathcal{T} = \{0, \dots, T-1\}$, with starts τ_t, ends τ_{t+1} and durations

$$\Delta_t = (\tau_{t+1} - \tau_t)/\text{one hour}. \tag{1}$$

The historical study uses quarter-hours, so $\Delta_t = 0.25$. A local day is constructed between successive midnights in `Europe/Rome`, then stored in UTC: T is 92, 96, or 100 for spring-DST, ordinary, or autumn-DST days. Repeated local hours are distinguished by their actual UTC offset/fold. No equation assumes that every day has exactly 24 elapsed hours.

A physical interval is not necessarily an independent market product. Before 1 October 2025, DA products are hourly; the MI auctions in the current 2025 history are already quarter-hourly. Physical battery controls are quarter-hourly at all four optimization stages. For the four equal-duration quarters of one native DA hour h, the implementation enforces

$$q_{\mathrm{DA},n,t} = q_{\mathrm{DA},n,t'} \quad (t, t' \in h). \tag{2}$$

These are equal MWh shares of one hourly order, whose total is $\sum_{t \in h} q_{\mathrm{DA},n,t}$. The hourly clearing price is repeated as an objective coefficient on those shares. The quarters do not become separate DA trades or separate observed prices. There is no corresponding date-based hourly equality on battery charging or discharging. Native product resolution and optional physical-control grouping are separate configuration fields.

The current study's configured local calendar is:

Auction	Gate	Results time used	Eligible delivery
DA	D−1 12:00	D−1 13:30	Whole day D
IDA1	D−1 15:00	D−1 15:30	Whole day D
IDA2	D−1 22:00	D−1 22:30	Whole day D
IDA3	D 10:00	D 10:30	D 12:00 onward

Submissions are planned five minutes before the gate in the supplied study. These times, date validity, result-time provenance and cancellations are explicit YAML settings. In particular, the historical DA publication time is a study assumption. Gate time and the start of IDA3 delivery are different events. Recorded publication timestamps can override configured times.

An auction a has an eligibility set $\mathcal{T}_a$. A documented/configured cancellation removes that auction from the active set. Missing data alone does not establish a cancellation: unexplained missing or incomplete curves fail validation. Absent auctions create no trade variables and no submission solve.

2.2 Notation and electrical basis

Symbol	Unit	Meaning
E	MWh	Nominal internal stored-energy capacity.
$E^{\min}, E^{\max}$	MWh	Operational energy limits derived from SOC percentages.
$\overline{P}$	MW AC	One symmetric battery charging/discharging rating.
$P^{\mathrm{imp}}, P^{\mathrm{exp}}$	MW AC	Directional shared POI ratings.

Symbol	Unit	Meaning
η_c, η_d	fraction	Complete AC-bus-to-storage and storage-to-AC efficiencies.
σ	fraction/hour	Configured hourly self-discharge fraction.
$V_{s,t}$	MW AC	Forecast available PV in scenario s.
$\widetilde{V}_t$	MW AC	Effective realized PV used for execution.
$c_{n,t}, d_{n,t}$	MW AC	Battery charge/discharge commands at information node n.
$u_{n,t}$	MW AC	PV curtailment.
$i_{s,t}$	MW AC	Uncurtailed PV injection into the common AC bus.
$g_{s,t}$	MW AC	Net grid exchange; positive is export, negative is import.
$e_{s,t}$	MWh	Internal battery energy at an interval boundary.
$q_{a,n,t}$	MWh	Signed auction adjustment; positive sale, negative purchase.
$Q_{s,t}$	MWh	Final cumulative contracted delivery position.
$\pi_{a,s,t}$	EUR/MWh	Auction-price coefficient for interval t.
w_s	fraction	Scenario probability, summing to one.

Power variables are AC at the common bus/POI. Battery energy is internal stored energy. Farm nameplate PV capacity is AC. For pvlib, the DC reference rating is derived as $P_{\mathrm{DC,STC}} = P_{\mathrm{AC,rated}}/\eta_{\mathrm{ref}}$, where the reference factor is a separate, fixed configured assumption (currently 0.96). Forecasts receive the DC-to-AC and transformer conversions and are then capped at the farm AC rating. Measured historical `power` is already AC kW and is converted once to MW; no PV or transformer efficiency is applied to it. BESS efficiency is not applied to measured PV or direct PV exports.

2.3 Forecast vintages and observability

A weather forecast cycle v has source time t_v and availability time A_v. Without explicit source publication times, the configured convention is $A_v = t_v + 4$ hours. At decision cutoff k, the rolling policy chooses the latest eligible cycle with $A_v \leq k$ and required remaining-horizon coverage. Weather cycles about six hours apart are successive forecast vintages, not ensemble members. The current dataset supplies one deterministic weather member.

Measured interval-average PV becomes available after interval end plus configured telemetry delay. A forecast for future intervals is not replaced by future realized power in the ordinary auction replay. If a decision occurs inside an unfinished quarter-hour, the next boundary state is projected through that interval with forecast PV and the same execution policy. Its final realized average cannot yet be observed. The hindsight counterfactuals deliberately have future PV and price information and are labeled as such.

Policy-treated observations carry evidence times. Plateau evidence can require later samples; two-sided interpolation requires its right endpoint and applicable mask evidence. A finalized corrected time series therefore differs from the as-of view available to a past decision. This remains relevant even when a study uses actual-price foresight.

3 PV power from weather forecasts

3.1 Irradiance and orientation

The input mappings declare forecast shortwave radiation as GHI and direct radiation as horizontal beam irradiance (BHI), with the latter an explicit working assumption. Hourly radiation means are end-labelled and redistributed uniformly over contained quarters. Other fields follow their declared instantaneous/accumulation semantics rather than the same blanket interpolation. Solar geometry is evaluated at each quarter-hour midpoint using configured coordinates and farm elevation.

For eligible irradiance data,

$$G_{\mathrm{DHI}} = G_{\mathrm{GHI}} - G_{\mathrm{BHI}}, \qquad G_{\mathrm{DNI}} = \frac{G_{\mathrm{BHI}}}{\cos\theta_z} \quad \text{away from the low-sun cutoff.} \tag{3}$$

The code uses pvlib's DNI helper with configured low-sun handling and a finite DNI cap or invalidation rule. BHI exceeding GHI is either invalidated or explicitly clipped and flagged. Optional Erbs decomposition can replace missing/unusable direct input under its declared conditions; it is not a general repair of contradictory source values.

Fixed arrays use metadata tilt and a declared source-to-pvlib azimuth conversion. Tracking IDs 1 and 3 use `pvlib.tracking.singleaxis` with configurable axis tilt, axis azimuth, rotation limit, backtracking and ground-coverage ratio. Type 3 currently requires the explicit front-side-only approximation; rear-side bifacial gain is omitted. The current geometry uses an acknowledged Milan location proxy rather than verified farm coordinates.

With panel tilt β, incidence angle α and ground albedo a, the isotropic transposition can be written

$$G_{\mathrm{POA}} = G_{\mathrm{DNI}}\,[\cos\alpha]_+ + G_{\mathrm{DHI}}\frac{1+\cos\beta}{2} + aG_{\mathrm{GHI}}\frac{1-\cos\beta}{2}. \tag{4}$$

3.2 Temperature, DC output and AC conversion

Wind speed is the magnitude of the configured horizontal wind components, $v = \sqrt{v_u^2 + v_v^2}$. The Faiman temperature model is

$$T_{\mathrm{mod}} = T_{\mathrm{air}} + \frac{G_{\mathrm{POA}}}{U_0 + U_1 v}. \tag{5}$$

The supplied wind height is explicitly acknowledged as a module-wind proxy. The AC farm rating is converted to a DC reference rating by dividing by the fixed, configurable assumption η_{ref}, currently 0.96. This assumption is independent of the downstream DC-to-AC coefficient. The resulting temperature and DC reference rating are passed to pvlib's PVWatts DC calculation:

$$P_{\mathrm{DC}} = \max\left\{0,\ \frac{P_{\mathrm{AC,rated}}}{\eta_{\mathrm{ref}}}\frac{G_{\mathrm{POA}}}{1000}\left[1+\gamma(T_{\mathrm{mod}} - T_{\mathrm{ref}})\right]\right\}. \tag{6}$$

DC power is scaled to AC by the configured explicit PV-output efficiency η_{PV}, or by an explicitly fitted DC-to-AC coefficient κ:

$$\widehat{P}_{\mathrm{AC}} = \begin{cases} \eta_{\mathrm{tr}}\eta_{\mathrm{PV}}P_{\mathrm{DC}}, & \text{explicit-efficiency route,} \\ \eta_{\mathrm{tr}}\kappa P_{\mathrm{DC}}, & \text{fitted-coefficient route.} \end{cases} \tag{7}$$

A configured AC cap, limited to no more than the farm AC rating, clips this output; the night mask and validity flags are then applied. Without a fitted override, the model uses the configured fixed output and separate transformer factors. An offline calibration pairs each quarter-hour's latest

available weather vintage with measured AC production y_i and modeled DC d_i. With $x_i = \eta_{\text{tr}} d_i$, it estimates the zero-intercept slope by minimizing mean squared error:

$$\hat{\kappa} = \arg\min_{\kappa>0} \frac{1}{n} \sum_i (y_i - \kappa x_i)^2 = \frac{\sum_i x_i y_i}{\sum_i x_i^2}. \tag{8}$$

Low-DC and invalid observations are excluded according to the saved study configuration. A chronological holdout compares the training-only fit with the configured explicit coefficient; the reported final coefficient is then refitted on all eligible pairs. This is a predictive aggregate scale, not an identified physical inverter efficiency. Temperature and wind affect production through temperature; irradiance and orientation affect POA. Cloud/dew-point fields retained in the store are not additional direct PV predictors here.

3.3 Chronological physical PV calibration

The `study-pv-calibration` workflow compares physical models separately for each farm. It constructs the weather view available at each auction's submission cutoff, including the configured six-hour UTC lag and decision buffer. IDA3 evaluation covers only local noon onwards. Targets are raw quarter-hour measured AC power; no transformer loss, interpolation or forecast substitution is applied to the targets. Missing/invalid targets are excluded by a common mask; genuine zeros remain. Satellite GHI diagnoses weather error without entering an operational prediction.

Let $\mathcal{T}_{da}$ be the eligible daylight targets for delivery day d and auction view a, and $f(X_{dat}; \theta)$ the shared physical predictor, including its AC cap and night policy. The fitted objective is

$$\min_{\theta \in \mathcal{B}} \frac{1}{|\mathcal{G}|} \sum_{(d,a)\in\mathcal{G}} \frac{1}{|\mathcal{T}_{da}|} \sum_{t\in\mathcal{T}_{da}} \rho\left(\frac{f(X_{dat};\theta) - P_{dt}^{\text{measured}}}{P_{\text{AC,rated}}} \right) + \lambda \sum_j \left(\frac{\theta_j - \theta_{0,j}}{u_j - l_j} \right)^2. \tag{9}$$

Here $[l_j, u_j]$ are configured bounds, θ_0 denotes the configured assumptions, and ρ is squared error or Huber loss. Equal day/auction weights prevent the shorter IDA3 horizon being underweighted. Scale-only fits optimize the clipped forward calculation; the earlier closed-form slope is retained only by the legacy single-coefficient command.

Candidate families fit κ alone, (κ, U_0, U_1), or additionally tracker GCR and maximum rotation for tracking types 1 and 3. The DC reference, γ, transformer factor, known fixed-array orientation and tracker-axis assumptions remain fixed. Front-side candidates compare isotropic, Hay–Davies and Perez–Driesse transposition. For farm 901, an additional infinite-sheds family uses explicit bifaciality and height/pitch assumptions, holding geometry fixed while fitting scale and thermal coefficients. Electrical effective POA includes the weighted rear irradiance, while Faiman uses incident front plus rear POA. One-at-a-time rear-assumption sensitivities are reported separately. AC power alone does not establish that fitted coefficients are true equipment properties; bound hits, optimizer stops and local Jacobian diagnostics are saved.

January is warmup and February–December 2025 are held-out delivery months. Monthly fits are issued at local midnight on the preceding calendar day, before the month's first DA submission. Only already available completed days enter training. The latest seven days select the history window (expanding or 90 days), transposition and regularization, after which the specification is refitted on the available history. A separate selected strategy chooses among families using that earlier validation alone. A final all-year fit is labeled retrospective/in-sample and never included in held-out scores. Paired accuracy intervals resample seven-day blocks; daily energy errors remain separate by auction because auction views overlap.

The `retrain-pv-calibration` command reuses this fitting procedure at a configured historical-availability timestamp. Immutable artifacts retain farm metadata, training/issue cutoffs, source checksums, software versions and fixed physical assumptions. The optional `pv.parameter_reg`

`istry` selects only the latest matching artifact issued by a forecast's decision cutoff, using the fixed baseline before the first fit. This interface is shared by ordinary PV forecasts and auction experiments and is disabled by default.

For ordinary experiments, the shared configuration now stores a final static parameter set under each farm's `pv` mapping. The model family is chosen by the full held-out comparison and its final all-2025 refit supplies the numbers. By the user's modeling convention, these inferred quantities are treated as prior farm metadata and apply at every experiment date, including January 2025. Their full-year fit scores are in-sample. This static assumption is distinct from the causally issued monthly benchmark; that study explicitly disables the farm mappings to preserve its original unfitted reference.

3.3.1 How the installed 2025 parameter sets were selected

The completed run is `outputs/pv_calibration/physical_pv_all_farms_2025/`. It contains 365 delivery days per farm, with January warmup and 334 February–December evaluation days. It completed 312 method/issue checkpoints across the four farms. A finite shared eligibility mask is frozen before model comparisons; the 12 missing unique target quarters per farm are not filled or scored. The selection has three separate levels:

1. For each method and monthly issue, fit on the earlier history and validate on its latest seven completed days. At least 14 earlier training days are required. Compare expanding and trailing 90-day history, admissible transposition models, and $\lambda \in \{0, 0.01\}$. Scale-only fits retain the reference transposition and use $\lambda = 0$. Rank candidates by daylight, AC-normalized RMSE with equal aggregate weight per day/auction view; then refit that candidate on its entire available history, including validation. The Huber transition is 0.05 of AC rating.

2. Evaluate the resulting monthly forecasts on the following delivery month. For the static defaults, select the method family with the smallest aggregate February–December daylight RMSE. This is a hindsight family selection; the separate monthly `selected` strategy uses only prior validation and can choose a different family each month.

3. For the selected family, use its final all-2025 refit. This final fit uses expanding history only; its transposition and regularization are chosen on the latest seven training days, and its physical parameters are then refitted on all 365 days. These are the values installed in `configs/standard.yaml`, with provenance artifacts in `configs/pv_parameters/2025/`. By the agreed convention they are applied as static prior metadata at every delivery date.

This is bounded nonlinear physical calibration, not a regression on arbitrary weather features. The scale-only problem uses bounded scalar minimization; thermal/tracker/bifacial fits use L-BFGS-B in unit-box parameter coordinates. The completed benchmark used the configured reference start plus one seeded random start, a budget of 180 function evaluations per start, and retained the best finite iterate, including the starting reference. The L^2 term regularizes deviations from the original physical assumptions; no L^1 penalty is used. All four installed final fits selected $\lambda = 0$.

Farm	Family/loss	Transposition	κ	U_0	U_1	GCR	Rotation
1	Thermal/Huber	Isotropic	0.531940	60	15	—	—
345	Thermal/Huber	Isotropic	0.462625	60	15	—	—
901	Tracker/MSE	Isotropic	0.889936	60	7.835820	0.638566	57.9135°
910	Tracker/MSE	Perez–Driesse	0.805011	60	15	0.566721	30°

Table 3: Installed all-2025 effective physical parameters, rounded for display. YAML retains full precision. The DC reference ratio remains 0.96 independently of κ; measured AC has no additional conversion loss.

The explicit rear-side candidate did not win farm 901's held-out comparison; its installed forecast therefore retains the calibrated front-side tracker approximation. The thermal parameters often reach their upper bounds and farm 910's rotation reaches the lower bound. The final farm 901 fit exhausted its configured numerical budget before reporting convergence; the best finite iterate is retained and flagged in its artifact. The other three final fits report convergence. Consequently these are the best predictive parameters found within the configured search, not certified global optima or measured equipment properties. These diagnostics are part of the delivered comparison.

Farm	Original RMSE (MW)	Monthly family RMSE (MW)	Reduction
1	0.189674	0.076846	59.5%
345	0.333161	0.144849	56.5%
901	0.632407	0.571855	9.6%
910	1.119012	1.032802	7.7%

Table 4: February–December held-out daylight RMSE. Each calibrated column uses that family's monthly historical fits, not the final all-year parameters. Family selection itself uses the complete held-out comparison.

3.3.2 Average, minimum and maximum errors against the previous model

For each matched quarter-hour define the absolute AC power error $e_{dat} = |\widehat{P}_{dat} - P_{dt}^{\text{measured}}|$. The mean is $\sum_{d,a,t} w_{dat} e_{dat}$, using equal aggregate day/auction weights; the minimum and maximum are the extrema across those same quarter-hour pairs. The following table compares the original fixed assumptions with the final static parameters over all of 2025. It uses daylight intervals to avoid improvements driven by nighttime zeros. It is an *in-sample* description of the global fit, distinct from the preceding held-out RMSE table.

	Mean absolute (MW)		Minimum (MW)		Maximum (MW)		Mean reduction
Farm	Before	After	Before	After	Before	After	
1	0.144662	0.054573	0	0	0.725960	0.483678	62.3%
345	0.247601	0.108362	0	0	1.203064	0.921906	56.2%
901	0.443153	0.390569	0	0	3.090024	3.106609	11.9%
910	0.717493	0.681918	0	0	5.815742	5.383478	5.0%

Table 5: Matched 2025 daylight absolute errors with the final static parameters. Minimum errors are zero even after excluding nighttime. Genuine zeros and suspected operating outages remain in the common target population.

Average improvement does not imply improvement of every interval or of the worst error: farm 901's maximum increases slightly. Residual weather errors, proxy location and operating anomalies remain. The machine-readable `csv/before_after_global_parameters.csv` and its JSON companion report both full-day and daylight mean/minimum/maximum errors, signed bias and RMSE in MW and percent of installed AC, separately by auction and in aggregate. They include a separately labeled monthly held-out-family comparison on its own common February–December population. Overlapping auction views are never interpreted as additive production or energy.

The four figures under `plots/trajectory_examples/` use the operational forecast function with these installed defaults. Their three common dates (24 January, 25 September and 6 September 2025) are chosen nearest the 20th, 50th and 80th percentiles of portfolio-normalized measured daily energy, independently of forecast residuals. Each panel shows unchanged raw measured AC, the DA/IDA1/IDA2 curves, the IDA3 noon-onward curve, the following midnight sample when available, and the installed AC rating as a fixed threshold. All 12 displayed farm/day examples

have no forecast or measured rating exceedance. Values, threshold checks, effective configuration and vintage provenance accompany the plots in CSV and JSON. These figures illustrate the static model on its fitting year rather than held-out forecasting skill.

Invalid required inputs remain invalid unless an explicit policy applies. Historical GME studies may use their configured night-zero forecast policy and a bounded fallback for missing realized-power intervals. Both the number and identity of substituted intervals are recorded. All counterfactuals use the same *effective* realized PV series, including these substitutions.

4 Optimization at an auction decision

Each residual solve uses only the remaining suffix of the delivery grid. In the equations below that suffix is reindexed from zero, with T denoting its final boundary; its original UTC product identifiers are retained for commitments.

4.1 Scenario tree and nonanticipativity

At decision k, let $\mathcal{S}_k$ be represented leaf scenarios and let $\mathcal{A}_k$ be the open active auctions. An information tree assigns an auction bid owner $n_a(s)$ to each auction/leaf and a control owner $n_t(s)$ to each interval/leaf. Curtailment may have a separate owner $m_t(s)$ under the selected execution-information convention.

Decisions are node-owned variables, not independent leaf decisions followed by a penalty for disagreement. For example,

$$c_{s,t} = c_{n_t(s),t}, \quad d_{s,t} = d_{n_t(s),t}, \quad q_{a,s,t} = q_{a,n_a(s),t}, \quad u_{s,t} = u_{m_t(s),t}. \tag{10}$$

Thus two scenarios with the same observable history share actionable decisions. A latent PV-member label is not available information by itself. Constraints apply to every represented leaf, including zero-probability leaves if supplied; probabilities only weight the objective. In the current one-scenario historical price oracle, the tree degenerates to deterministic planning with known future prices and a forecast PV trajectory.

4.2 Physical variables and balances

For every scenario and remaining interval:

$$i_{s,t} + u_{s,t} = V_{s,t}, \qquad i_{s,t} \geq 0, \qquad 0 \leq u_{s,t} \leq V_{s,t}, \tag{11}$$

$$g_{s,t} = i_{s,t} + d_{s,t} - c_{s,t}, \tag{12}$$

$$-P^{\text{imp}} \leq g_{s,t} \leq P^{\text{exp}}. \tag{13}$$

Curtailment appears only in the PV balance. It cannot erase bought grid energy or battery charging energy. A negative grid exchange must be absorbed by net battery charging since there is no local demand. The model does not explicitly assign charging power to its PV/grid source; that partition is an accounting convention used in post-processing.

With $z_{n,t} \in \{0, 1\}$ denoting charging mode,

$$0 \leq c_{n,t} \leq \overline{P} z_{n,t}, \qquad 0 \leq d_{n,t} \leq \overline{P}(1 - z_{n,t}). \tag{14}$$

This excludes simultaneous charge and discharge, uses one symmetric AC rating, and introduces no separate battery-inverter power limit.

SOC percentage parameters $s_{\min}, s_{\max}$ define

$$E^{\min} = E\frac{s_{\min}}{100}, \qquad E^{\max} = E\frac{s_{\max}}{100}, \qquad E^{\min} \le e_{s,t} \le E^{\max}. \tag{15}$$

For $\rho_t = (1-\sigma)^{\Delta_t}$, the transition is

$$e_{s,t+1} = \rho_t e_{s,t} + \eta_c c_{s,t}\Delta_t - \frac{d_{s,t}\Delta_t}{\eta_d}. \tag{16}$$

One AC MWh charged stores η_c MWh. Delivering one AC MWh consumes $1/\eta_d$ internal MWh. In combined mode the two configured efficiencies each cover their entire directional path. In converter-and-cell mode, $\eta_c = \eta_{c,\text{conv}}\eta_{c,\text{cell}}$ and $\eta_d = \eta_{d,\text{conv}}\eta_{d,\text{cell}}$. The alternatives cannot be mixed. The round-trip factor is $\eta_c\eta_d$. Self-discharge acts on stored energy, not on a power variable, and is separate from cycle degradation cost.

The initial boundary is fixed to measured/carried/projected state e^k:

$$e_{s,0} = e^k. \tag{17}$$

If the terminal target is configured, also impose

$$e_{s,T} = e^{\text{target}}. \tag{18}$$

In the GME adapter, initial/terminal percentages are converted using nominal capacity. A `null` terminal SOC omits (18). Prepared and synthetic boundary-state contracts use explicitly named MWh values. Neither version adds a terminal energy salvage value to the objective.

4.3 Auction quantities and physical delivery

Let Q_t^{fix} contain all previously submitted commitments, priced or still pending. The new final position on leaf s is

$$Q_{s,t} = Q_t^{\text{fix}} + \sum_{a\in\mathcal{A}_k:\, t\in\mathcal{T}_a} q_{a,n_a(s),t}. \tag{19}$$

An ineligible auction/interval has no trade variable. The model allows signed sales/purchases; they are adjustments to prior positions, not gross generation. For every eligible trade,

$$-P^{\text{imp}}\Delta_t \le q_{a,n,t} \le P^{\text{exp}}\Delta_t. \tag{20}$$

For every open auction a in chronological order, define the cumulative planned position after that auction as

$$C_{s,t}^{a} = Q_t^{\text{fix}} + \sum_{b\in\mathcal{A}_k:\, b\le a,\, t\in\mathcal{T}_b} q_{b,n_b(s),t}. \tag{21}$$

It satisfies

$$-P^{\text{imp}}\Delta_t \le C_{s,t}^{a} \le P^{\text{exp}}\Delta_t, \qquad -P^{\text{imp}}\Delta_t \le Q_{s,t} \le P^{\text{exp}}\Delta_t. \tag{22}$$

There is no independent commercial volume cap. The directional POI ratings apply to individual auction adjustments, cumulative positions after each auction, and physical delivery. An offsetting later trade is allowed, subject to its own POI-derived limit and the resulting position limit. Thus the model permits bounded cross-auction spread trading even without a battery; its revenues can be present in both asset counterfactuals.

With `trading.optimize_imbalances: false`, which is the current farm-910 policy, planned delivery is exact in every interval:

$$g_{s,t}\Delta_t = Q_{s,t}. \tag{23}$$

No positive/negative imbalance variables or imbalance direction binary are created. Forecast deviations are not an additional dispatch resource. An optional configured product-group delivery equality can additionally impose $\sum_{t\in h}(g_{s,t}\Delta_t - Q_{s,t}) = 0$. It is redundant with (23) and does not replace it in the disabled-imbalance mode.

4.4 Current objective and sunk cash flows

For throughput degradation with coefficient c_{thr} and an optional import fee f, the current residual objective is

$$\max J_k = \sum_{s \in \mathcal{S}_k} w_s \left[\sum_{a \in \mathcal{A}_k} \sum_{t \in \mathcal{T}_a} \pi_{a,s,t} q_{a,n_a(s),t} - c_{\text{thr}} \sum_t (c_{s,t} + d_{s,t}) \Delta_t - f \sum_t [-g_{s,t}]_+ \Delta_t - C_s^{\text{trading}} \right]. \tag{24}$$

The implementation uses an exact import-power linearization with a direction binary only when f is nonzero. The farm-910 study sets $f = 0$.

4.5 Market transaction fees and operator monthly volume

The `transaction_fees.mode` is `none`, `flat`, or `gme_real` (the shared YAML default). Legacy programmatic inputs without a fee section remain fee-free. Let $m(a)$ be the calendar month of auction submission in market local time, V_m known already executed operator volume, and $G_{m,s} = \sum_{a:m(a)=m,t} |q_{a,n_a(s),t}|$ new gross MWh. Purchases and sales count separately even when final delivery is zero. Quarter-hour energy shares of an hourly product are summed once, without another duration factor. For a cumulative progressive tariff F_m,

$$C_s^{\text{trading}} = \sum_m \left[F_m(V_m + G_{m,s}) - F_m(V_m) \right]. \tag{25}$$

In flat mode this is the configured EUR/MWh coefficient times gross MWh. The published GME MPE marginal rates are 0 on the first 20,000 MWh/month, 0.04 up to 1,000,000, 0.03 up to 10,000,000, and 0.02 thereafter. These bands are progressive and not a retroactive rate on the monthly total.

Fee-free horizons introduce no auxiliary variables. Positive single-band costs use linear absolute-value epigraphs. Across reachable breakpoints, exact Gurobi ABS general constraints and a PWL equality represent both gross volume and the nonconvex tariff; a convex epigraph alone is not used for regressive bands. Absolute-value variables retain auction-node ownership and native-product sharing. Costs are evaluated for each leaf and weighted by w_s.

The fee state includes pending guaranteed transactions immediately at submission. Publication adds no second fee. Executed gross volume remains in V_m during residual solves, while incurred EUR are sunk. The state carries an availability clock and declared provenance; nonzero future executed monthly volume is rejected. Months use transaction dates, so pre-DA/IDA1/IDA2 and IDA3 for one delivery day may belong to different months. A new month starts at zero. The saved final fee state can initialize a later chronological study. Operator scope must explicitly distinguish a standalone plant from a portfolio with exogenous starting volume. Unknown future external activity is not invented: costs are exact conditional on the supplied initial state, not a complete operator-level monthly forecast.

Tariffs have explicit nonoverlapping complete-month validity windows, version, source, thresholds, marginal rates and an applicability status. The current GME tariff and the 2006 annual report show the same structure, the latter effective in 2007. A dated confirmation of 2025 applicability was not established; the shared December 2024–December 2026 tariff marks this explicitly as `assumed_unchanged`. Unsupported dates fail. The shared initial state assumes a standalone simulated operator at local midnight on 1 December 2024, with zero December volume, zero earlier fees and no external portfolio activity. Subsequent months begin at zero and continuous execution carries actual submitted volume. Known operator-level volume should replace this assumption when the plant belongs to a larger portfolio. Enabling this tariff can still produce zero variable fees below the monthly exemption. Example YAMLs are `configs/full_experiments/s`

ynthetic_fees_none.yaml, configs/full_experiments/synthetic_fees_flat.yaml, and configs/full_experiments/synthetic_fees_gme.yaml. Membership, taxes, ACER record charges and separate aggregator commissions remain excluded; physical grid fees and degradation remain separate cost components.

Past auction cash flow and already incurred degradation are sunk at residual reoptimization. Priced commitments remain in the external ledger, and pending unpriced submitted quantities enter Q^{fix} without adding an optimizable cash-flow term. Their future clearing value is constant with respect to all remaining decisions and is excluded from J_k. This avoids double counting. The initial pre-DA solve has no prior commitments; its objective spans the modeled whole-day opportunity set. Residual objective values therefore should not be summed across auction solves to compute daily profit.

There is no explicit curtailment penalty, risk measure, capital cost, fixed O&M, tax, brokerage fee, ramp constraint, minimum dwell time, or calendar-aging charge in (24). Curtailment can affect revenue through delivery choices, but is not independently priced. The horizon can exploit inventory when its terminal target is free; this matters when comparing asset value.

4.6 What is free at each of the four decisions

Decision	New free auction quantities	Previously submitted fixed quantities
DA	DA, IDA1, IDA2, IDA3 plans	None
IDA1	IDA1, IDA2, IDA3 plans	DA
IDA2	IDA2, IDA3 plans	DA and IDA1
IDA3	IDA3 eligible products	DA, IDA1 and IDA2

At each decision, remaining-horizon charge/discharge modes, curtailment, grid powers and energy states are free subject to their balances, boundary state, nonanticipativity, product rules and terminal condition. Executed past intervals are absent from the residual grid. Future-auction quantities are plans until their own submission. Only the current auction order is committed from that solve. Before noon, IDA3 cannot change orders for the morning; physical control commands from the previous plan also remain in force. All IDA3 decision variables cover only the 48 quarter-hours from 12:00 to midnight: charge, discharge, mode, curtailment, grid power and trades, with 49 SOC boundaries. Gate closure determines the information cutoff, not the delivery horizon.

More generally, let c_k be the submission cutoff and j_k the first grid boundary no earlier than both c_k and the current auction's first eligible delivery. The solver starts at j_k. From the latest observed physical state, the replay projects the locked prefix using the previously requested commands $\bar{u}_i$, fixed contracted grid targets Q_i/Δ_i, and the latest cutoff-available PV forecast $\widehat{P}_i^{(c_k)}$ through the same physical controller Φ used for delivery:

$$\widehat{x}_{i+1|c_k} = \Phi\Big(\widehat{x}_{i|c_k}, \bar{u}_i, Q_i/\Delta_i, \widehat{P}_i^{(c_k)}\Big), \qquad i < j_k. \tag{26}$$

Here x includes energy and degradation state. The resulting $\widehat{E}_{j_k|c_k}$ is the fixed initial energy of the residual MILP. Forecast publication eligibility is still checked at c_k; projecting to noon does not permit a forecast released after the 10:00 gate. No future measured PV is used in this projection. Actual delivery continues with the retained commands and realized PV, so its noon state may differ from the projected state. Grid feedback and protection enforce the actual SOC and power bounds. If IDA3 is the first participating auction, zero-position feedback continues until noon.

5 Battery degradation models

5.1 Current AC throughput approximation

For one trajectory,

$$H = \sum_t (c_t + d_t)\Delta_t, \qquad D_{\text{thr}} = c_{\text{thr}} H. \tag{27}$$

Both directional AC throughputs are charged once. The YAML can directly supply c_{thr} or derive it from replacement economics and reference cycle life. Let p_{USD} be pack USD/kWh, x be USD per EUR, p_{extra} additional EUR/kWh replacement cost, s_R salvage fraction, N reference cycle life and d_{ref} reference internal depth of discharge. Then

$$R = 1000E\left(\frac{p_{\text{USD}}}{x} + p_{\text{extra}}\right)(1 - s_R), \tag{28}$$

$$H_{\text{life}} = N d_{\text{ref}} E\left(\frac{1}{\eta_c} + \eta_d\right), \tag{29}$$

$$c_{\text{thr}} = \frac{R}{H_{\text{life}}}. \tag{30}$$

One reference internal cycle charges $d_{\text{ref}}E/\eta_c$ AC MWh and discharges $d_{\text{ref}}E\eta_d$ AC MWh, explaining the normalization. There is no additional factor of two or half applied to (30). The configured end-of-life capacity fraction is descriptive metadata, not an extra multiplier or a dynamic capacity update.

The current inherited farm-910 YAML uses a historical generic pack-price proxy, a dated FX conversion, and an illustrative reference cycle life. Those are study inputs with source metadata, not an estimate of a specific installed BESS cost. The resolved R and c_{thr} are saved. The cost is financial wear; it does not reduce E, $\overline{P}$, or efficiencies as the simulation proceeds.

5.2 Optional segmented depth-dependent aging

The optional `xu_segmented` mode follows the configured segmented cycle-depth construction described in the repository's Xu notes and the author paper [1]. It is distinct from the current throughput study. For a configurable convex stress curve $\Phi(d) = \alpha d^\beta$, $\beta \geq 1$, partition nominal internal capacity into J equal segments of size E/J. The marginal AC discharge cost of segment $j = 1, \ldots, J$ is

$$m_j = \frac{RJ}{\eta_d E}\left[\Phi(j/J) - \Phi((j-1)/J)\right]. \tag{31}$$

It is nondecreasing with depth. The degradation charge is

$$D_{\text{Xu}} = \sum_t \sum_{j=1}^{J} m_j d_{j,t} \Delta_t. \tag{32}$$

Only discharging is charged under this convention. A discharging half-cycle incurs the convention's depth cost, while a charging half-cycle does not; do not add an extra factor of one half. Nonzero throughput cost is rejected in this mode.

Segment flows and energies obey

$$\sum_j c_{j,t} = c_t, \qquad \sum_j d_{j,t} = d_t, \qquad \sum_j h_{j,t} = e_t, \tag{33}$$

$$h_{j,t+1} = h_{j,t} + \Delta_t(\eta_c c_{j,t} - d_{j,t}/\eta_d), \qquad 0 \leq h_{j,t} \leq E/J. \tag{34}$$

Segment energies represent cycle-depth memory, not fixed absolute SOC buckets. Charging and discharging both use shallow available segments first. For each $j < J$, the implementation adds priorities $v^c_{j,t}, v^d_{j,t} \in \{0,1\}$:

$$\sum_{\ell>j} c_{\ell,t} \le \overline{P} v^c_{j,t}, \qquad h_{j,t+1} \ge (E/J) v^c_{j,t}, \tag{35}$$

$$\sum_{\ell>j} d_{\ell,t} \le \overline{P} v^d_{j,t}, \qquad h_{j,t+1} \le (E/J)(1 - v^d_{j,t}). \tag{36}$$

Using a deeper segment therefore requires the shallower segment to end full when charging or empty when discharging. This chooses a deterministic shallow- first representative, including otherwise indifferent charging allocations. Segment actions have the same information ownership as aggregate battery power.

At a declared fresh study start, initialize

$$h_{j,0} = \mathrm{clip}\left(e_0 - (j-1)E/J,\ 0,\ E/J\right). \tag{37}$$

Afterwards the segment vector must be carried with actual execution; an arbitrary residual solve cannot reconstruct it from aggregate SOC. Previous throughput or aging requires explicit memory. Xu mode currently rejects nonzero self-discharge. It remains a finite-segment approximation and omits calendar, temperature and current-rate aging, capacity-fade feedback and unknown pre-study cycle history.

6 Submission, publication and reoptimization workflow

Before each active auction submission, the GME replay refreshes the latest available PV forecast over unexecuted intervals, projects a boundary if necessary, and solves the residual planning problem. Previously submitted quantities, including still-unpriced ones, are fixed. At submission, only the current auction's node-shared action is extracted. Later planned auction quantities are not booked and can change.

At publication, known prices are updated, and actual contractual revenue $\pi^{\mathrm{real}}_{a,t} q^{\mathrm{sub}}_{a,t}$ is entered exactly once in the ledger. Publication does not reopen a market. The current GME experiment forces `reoptimize_physical_controls=false`; it uses four auction solves for four active auctions, not an extra set of publication solves or 96 physical optimizations. Physical feedback is algebraic. Synthetic replay can separately enable interval-control reoptimization, so its solve count need not be four.

At the next auction, closed commitments are carried, remaining open opportunities are reconsidered, and past executed state/aging remain fixed. In the actual-price oracle all future auction prices are already supplied as a benchmark, so price publication does not remove price uncertainty there. In a forecast-based route, known results condition subsequent price distributions without changing the commitments that produced them.

7 Physical feedback with realized PV

7.1 Feasible interval power range

The current policy is `track_committed_grid`. For interval t its target is

$$g^\star_t = Q^{\mathrm{committed}}_t / \Delta_t. \tag{38}$$

The controller uses effective realized interval-average PV $\widetilde{V}_t$; this is an idealized interval simulator, not a subsecond telemetry/control model. It does not disclose this final average to the earlier auction optimizer. Given actual initial energy e_t and retained energy $\bar{e}_t = \rho_t e_t$, the feasible charging/discharging maxima are

$$C_t = \min\left\{\overline{P}, \frac{[E^{\max} - \bar{e}_t]_+}{\eta_c \Delta_t}\right\}, \tag{39}$$

$$D_t = \min\left\{\overline{P}, \frac{\eta_d \left[\bar{e}_t - E^{\min}\right]_+}{\Delta_t}\right\}. \tag{40}$$

First clip the planned charge/discharge commands to these maxima, obtaining c_t^0, d_t^0. Simultaneous positive requested charging and discharging after protection is rejected. Initially curtail only available PV: $u_t^0 = \min\{\widetilde{V}_t, \left[u_t^{\text{plan}}\right]_+\}$.

7.2 Grid-tracking policy

Release scheduled curtailment if it can help meet a deficit:

$$g_t^0 = \widetilde{V}_t - u_t^0 + d_t^0 - c_t^0, \tag{41}$$

$$\bar{u}_t = u_t^0 - \min\{u_t^0, \left[g_t^\star - g_t^0\right]_+\}. \tag{42}$$

The desired net battery dispatch and its feasible projection are

$$b_t^{\text{req}} = g_t^\star - \widetilde{V}_t + \bar{u}_t, \tag{43}$$

$$b_t = \text{clip}(b_t^{\text{req}}, -C_t, D_t), \tag{44}$$

$$c_t = [-b_t]_+, \qquad d_t = [b_t]_+. \tag{45}$$

Finally curtail any remaining surplus that battery adjustment cannot absorb:

$$u_t = \bar{u}_t + \min\{\widetilde{V}_t - \bar{u}_t, \left[\widetilde{V}_t - \bar{u}_t + b_t - g_t^\star\right]_+\}. \tag{46}$$

Thus excess PV first decreases discharge, then increases charging as needed and feasible, then is curtailed. A PV shortfall first releases planned curtailment, then reduces charging or increases discharge within the actual SOC box. The remaining tracking error is preserved if battery power/energy cannot supply it.

POI protection follows. If net import exceeds its rating, charging is reduced. If export exceeds its rating, discharge is reduced first and remaining excess PV is curtailed. Grid power is recomputed consistently with the adjusted actions. Actual energy then advances using (16) with these executed powers, and bounds are checked before accepting the new state. Both throughput/Xu memory and cumulative financial wear are advanced from executed actions, not from the old auction plan. Protection is mandatory under either execution policy; the alternative `follow_battery_plan` skips grid-tracking adjustment.

There is no controller constraint forcing the actual terminal energy to the auction plan's target. Forecast errors and grid tracking can cause an ending-SOC difference even if every optimization imposed (18). This is visible in the current terminal-50% study and is not equivalent to guaranteeing 50% realized SOC. Any hard actual-terminal policy would be an additional controller design and could create different imbalances/cash flows.

7.3 PV source allocation for reporting

For the executed trajectory, use the explicit PV-first accounting convention

$$P_t^{\mathrm{PV}\to\mathrm{B}} = \min\{\widetilde{V}_t - u_t, c_t\}, \tag{47}$$

$$P_t^{\mathrm{PV}\to\mathrm{G}} = \widetilde{V}_t - u_t - P_t^{\mathrm{PV}\to\mathrm{B}}, \tag{48}$$

$$P_t^{\mathrm{PV}\to\mathrm{curt}} = u_t. \tag{49}$$

They sum to $\widetilde{V}_t$ in every interval. Grid-supplied charging and battery discharge are not counted as PV production. These are post-processing flows; there are no separate source-attribution decision variables in the MILP.

8 Settlement and expected-versus-realized accounting

Actual net deviation is retained at the native quarter-hour:

$$\delta_t^{\mathrm{real}} = g_t^{\mathrm{real}}\Delta_t - Q_t^{\mathrm{committed}}. \tag{50}$$

Positive means more net export than promised. With the selected single Terna NORD settlement price λ_t, the modeled cash flow is

$$S_t = \lambda_t \delta_t^{\mathrm{real}}. \tag{51}$$

A positive price therefore does not mean every imbalance is penalized: positive excess export is credited and a negative deviation is debited. Negative prices are allowed and reverse these cash-flow signs. The general settlement helper also accepts two directional curves:

$$S_t = \lambda_t^{+}\left[\delta_t^{\mathrm{real}}\right]_+ - \lambda_t^{-}\left[-\delta_t^{\mathrm{real}}\right]_+. \tag{52}$$

This documents the implemented study accounting, not every possible Italian participant-specific settlement item or regulatory regime. The workbook adapter selects the declared NORD macrozone and the imbalance-price column, retaining other imported fields as provenance. Product starts and ends must match the execution grid exactly. Prices are neither hourly-averaged nor filled across gaps; missing enabled settlement prices are an error.

Actual daily auction cash flow and settled operating profit are

$$A^{\mathrm{real}} = \sum_{a,t} \pi_{a,t}^{\mathrm{real}} q_{a,t}^{\mathrm{sub}}, \tag{53}$$

$$J^{\mathrm{real}} = A^{\mathrm{real}} + \sum_t S_t - D^{\mathrm{executed}} - f\sum_t \left[-g_t^{\mathrm{real}}\right]_+ \Delta_t - C^{\mathrm{trading,real}}. \tag{54}$$

A model deterrence penalty is not an actual settlement charge and is not subtracted a second time by actual post-processing. If settlement is unavailable, profit before settlement is reported and fully settled profit remains unknown. Only actually submitted guaranteed transactions enter $C^{\mathrm{trading,real}}$; future plans that are never submitted contribute no realized fees. The ledger retains transaction-month volume, per-trade charges and tariff versions. An independent Python calculation reconciles those charges and the final monthly state. Expected and realized transaction fees are reported separately from gross auction revenue. In `realized_profit_breakdown.png`, trading-fee EUR and gross MWh occupy the former import-fee slot; nonzero physical grid fees appear as their own additional bar.

Let J^{init} and A^{init} be the initial auction solve's expected profit and auction revenue. Saved mismatches are

$$\Delta J = J^{\mathrm{real}} - J^{\mathrm{init}}, \qquad \Delta A = A^{\mathrm{real}} - A^{\mathrm{init}}, \qquad r_J = 100\frac{\Delta J}{\left|J^{\mathrm{init}}\right|}. \tag{55}$$

The percentage is undefined/`null` when its expectation is zero; unsettled profit comparisons remain unknown. Differences include later replanning and physical execution effects, not only price-forecast error.

The replay also evaluates each stored auction snapshot as an open-loop diagnostic against realized PV and the same settlement prices. Its planned battery commands are protected physically. Deviations are measured against that snapshot's planned grid schedule (which can include assumed future trades) or the final contract, as labeled in the recorded comparisons. It is not another auction optimization. Because IDA3 covers a shorter residual horizon, a common-horizon comparison is also saved; unaligned whole-horizon totals should not be compared as if their delivery intervals were identical.

9 Value of reoptimization through frozen-plan replay

The optional frozen-plan comparison asks what would happen if planning stopped after one auction decision but physical feedback continued. This differs from the older open-loop imbalance diagnostic, which follows planned battery commands without grid-tracking feedback and does not reconstruct an alternative complete auction ledger and net profit.

Let k be an active auction other than the final active auction. At its cutoff, retain the entire deterministic future-trade vector from that solve. Define

$$q_{a,t}^{[k]} = \begin{cases} q_{a,t}^{\text{submitted,ref}}, & a \le k, \\ q_{a,t}^{\text{plan},k}, & a > k, \end{cases} \qquad Q_t^{[k]} = \sum_a q_{a,t}^{[k]}. \tag{56}$$

The already submitted trades are fixed even if their prices are unpublished. Future orders are the frozen snapshot's orders, not the later realized orders from the reoptimized strategy. They are submitted at their original gates and priced at actual clearing prices. Their eligibility must ensure submission precedes the relevant delivery interval. Thus their cumulative position is a valid delivery-time tracking target without prematurely revealing new prices or forecasts to the frozen plan.

Executed history before the frozen residual boundary remains identical to the reference trajectory. The suffix starts from the actual carried energy, throughput, incurred wear and, for Xu mode, segment-memory vector. It is not reset to the snapshot's forecast SOC, including when the cutoff projects a not-yet-observed boundary. At physical delivery that boundary's actual state is carried forward without changing the frozen decisions. The same configured execution policy uses the frozen battery/curtailment commands, target $Q_t^{[k]}/\Delta_t$, and the same effective realized PV:

$$(c_t^{[k]}, d_t^{[k]}, u_t^{[k]}, e_{t+1}^{[k]}) = \mathcal{F}(e_t^{[k]}, \widetilde{V}_t, Q_t^{[k]}/\Delta_t, c_t^{\text{plan,k}}, d_t^{\text{plan,k}}, u_t^{\text{plan,k}}). \tag{57}$$

Here $\mathcal{F}$ is the grid-feedback/protection policy described above. The current study uses grid tracking; the alternative execution policy, if selected, is shared with the reference. No later optimization or weather-forecast update is applied. The algebraic replay adds zero optimization solves.

Using the same signed settlement convention and actual dated transaction-fee accounting as the reference, the full-day frozen profit is

$$J_{\text{frozen}}^{[k]} = \sum_{a,t} \pi_{a,t}^{\text{real}} q_{a,t}^{[k]} + \sum_t S_t^{[k]} - D^{[k]} - C_{\text{import}}^{[k]} - C_{\text{trading}}^{[k]}. \tag{58}$$

Previously executed prefix costs are preserved. Future wear is recomputed from the counterfactual's executed commands; gross fees are recomputed from its own trade ledger and the same initial operator-volume state, including month transitions and tariff bands. Reference future trades or fees are never borrowed after the freeze point.

The value KPI is

$$G_k = J^{\text{real}}_{\text{fully reoptimized}} - J^{[k]}_{\text{frozen}}. \tag{59}$$

Both sides cover the entire delivery day, so differing residual lengths do not create incomparable profit totals. Positive values favor later reoptimization; negative values are retained. Each G_k measures all later replanning after k: the values are not additive and need not be monotonic. This comparison includes new forecasts, revised orders and state-aware planning rather than isolating a causal effect of forecast accuracy. Final stored-energy differences are reported; no terminal salvage price is assumed. If actual settlement is unavailable, settled profit and G_k remain unknown.

The independent YAML section is `reoptimization_baseline`. Its options are `enabled`, `freeze_after_auctions`, and `plots`. An absent section disables the feature; the comparison study enables it. A null auction selection evaluates every active pre-final auction, while a list chooses a subset. The last active auction is always excluded, cancelled auctions produce no baseline, and an empty selection produces no comparisons. The plot switch also respects the global plot setting. Snapshot replay requires deterministic plans; it rejects using one selected leaf of a contingent plan as if all of its future branches were already known.

The simulator is `src/simulation/frozen_plans.py`. Saved dispatch-plan rows include every `planned_trade_<auction>_mwh` and their planned net position. Full-day profits, gains and states are saved in the frozen-plan summary CSV/JSON, with separate actual counterfactual dispatch and trade ledgers. Individual gains also enter the scalar KPI table, and `reoptimization_value.png` compares frozen profits with the reference. The isolated January configuration is `configs/full_experiments/gme_oracle_2025-01-01_farm910_reoptimization_value.yaml`.

10 Perfect-information schedules and the PV-only baseline

10.1 PV+BESS versus the feedback controller

After replay, one additional full-day problem uses $V_t = \widetilde{V}_t$ and all actual auction-price curves, with a single scenario and all auction orders jointly free subject to their native products, eligibility, POI and physical constraints. It retains the BESS ratings, efficiencies, SOC box, degradation, initial energy and configured terminal target. Call its achieved solver objective $J^{\text{PI}}_{\text{PV+BESS}}$.

This is deliberately a hindsight plan rather than a causally executable policy. It is compared with the feedback controller's actual dispatch, curtailment and SOC, and its objective is saved as the solver's `ObjVal`. Four missing measurement intervals, if permitted by the selected run, are handled through that run's explicit fallback; diagnostics record their count and identities. The schedule figure shows both trajectories and their battery-power difference through the final day boundary. It does not assert that feedback should reproduce the hindsight trajectory or terminal energy.

10.2 PV-only asset counterfactual

When `experiment_outputs.pv_only_baseline: true`, a second independent solve uses the same optimizer builder with

$$E = 0, \quad \overline{P} = 0, \quad E^{\min} = E^{\max} = 0, \quad e_0 = e_T = 0, \quad D = 0. \tag{60}$$

Battery powers and energy variables are consequently fixed to zero; the code reuses one physical/trading formulation instead of implementing a second, potentially inconsistent PV-only model. Unused variables can be removed by solver presolve. PV injection/curtailment and auction trades remain free subject to the same effective realized PV, native products, active auctions and POI. Call its achieved objective $J^{\text{PI}}_{\text{PV}}$.

The incremental operating-value KPI is

$$\Delta J^{\mathrm{PI}}_{\mathrm{BESS}} = J^{\mathrm{PI}}_{\mathrm{PV+BESS}} - J^{\mathrm{PI}}_{\mathrm{PV}}. \tag{61}$$

The BESS side includes its own modeled degradation; the PV-only side has no battery wear. Both sides retain permitted cross-auction spread trading, so market-trading revenue common to both is not credited to battery ownership. The KPI does not subtract realized feedback wear from a different ideal schedule. Each counterfactual reports its own objective components. Directly subtracting a PV-only hindsight objective from the realized PV+BESS feedback profit would mix information and execution policies and is not the asset-value KPI.

This comparison is one-day operating value before capital expenditure, fixed operating costs, financing and taxes. A free terminal target can also use initial stored energy or leave different inventory; the selected 50% initial/terminal hindsight comparison avoids that particular inventory difference for the BESS model. Realized feedback may still finish at a different SOC. No residual-energy valuation, annualization, payback, NPV or investment optimization is implemented. Those need a representative horizon and additional cost/cash-flow assumptions.

Setting the switch to `false` skips the PV-only solve and its model, JSON, CSV and figure. Incremental-value metrics become `null` and an explicit baseline-enabled KPI is false. The perfect-information PV+BESS-versus-feedback comparison still runs. Four active auctions therefore produce four replay solves plus one PV+BESS hindsight solve and, when enabled, one PV-only solve. Counterfactual solve diagnostics are separate from replay solve counts.

11 Configurable participation and matched-day profile comparison

Participation is a choice of a nonempty subset P of the chronologically ordered auction set $\mathcal{A} = \{\mathrm{DA}, \mathrm{IDA1}, \mathrm{IDA2}, \mathrm{IDA3}\}$. The single-study configuration is `gme_oracle_experiment.participating_auctions`. Omission or null selects every configured auction dataset. Gate order remains unchanged; unselected auctions have no price-data requirement, trade variables, submission event/solve or fees. Their orders are absent rather than independent variables fixed after a solve. The active set on day d also removes documented/configured cancellations. If every selected auction is documented/configured as cancelled, the profile retains its full delivery day with zero orders, auction/control solves and auction fees. The zero-position feedback policy stores PV within the SOC box and curtails remaining surplus; executed wear and settlement remain in actual profit. There is no initial auction objective (its expectation KPIs are null), and no auction plan enforces the configured terminal target on that day. Missing or partial price data without cancellation evidence still fail.

Four independently optional auctions produce

$$|\{P \subseteq \mathcal{A} : P \neq \varnothing\}| = 2^4 - 1 = 15. \tag{62}$$

If DA is mandatory and the three IDA auctions remain optional, there are instead

$$|\{\{\mathrm{DA}\} \cup S : S \subseteq \{\mathrm{IDA1}, \mathrm{IDA2}, \mathrm{IDA3}\}\}| = 2^3 = 8. \tag{63}$$

Both families and explicitly listed profiles are supported. For example, $P = \{\mathrm{IDA1}, \mathrm{IDA2}\}$ and $P = \{\mathrm{IDA3}\}$ are valid without DA. All physical, native-product, degradation and settlement constraints continue to apply to the same full delivery day.

If the first selected gate occurs after delivery begins, its prefix uses zero contracted grid delivery and grid feedback with zero planned battery commands. The controller may store PV within the energy box and curtail remaining surplus; it does not earn uncontracted export revenue in this prefix. Actual energy and wear carry to the first selected solve, including the existing unfinished-interval state projection. No extra auction optimization is inserted before the first selected decision.

Prefix executed/projected costs are subtracted from the first residual objective when forming an initial whole-day expected profit; the residual model objective is separately retained. Actual whole-day profit already includes all executed prefix costs through the ordinary financial accounting.

11.1 Sequential current-auction-only baseline

At the decision for auction a, integrated planning includes the open auctions $\mathcal{A}_a$. The configurable sequential baseline replaces this set by $\{a\}$. Let $Q_t^{<a}$ denote all previously submitted, fixed quantities, including unpriced commitments. Its only free market quantity is $q_{a,t}$:

$$Q_t = Q_t^{<a} + q_{a,t}, \qquad \max \sum_t \lambda_{a,t} q_{a,t} - C_{\text{deg}} - C_{\text{import}} - C_{\text{fees}}. \tag{64}$$

The existing physical balances link Q_t to grid delivery. The full remaining delivery horizon retains charge/discharge, curtailment, SOC transitions, native product equalities, POI limits and the configured terminal target. No later auction trade or revenue appears. Sunk earlier cash flow does not affect the choice. Each next gate receives updated PV, actual carried battery state and immutable prior submissions. Under the price-oracle convention, the current auction's actual clearing curve is supplied to that solve.

`gme_oracle_experiment.auction_planning_mode` selects `integrated` (default) or `current_auction_only`. `participation_study.sequential_all_auctions_baseline` optionally adds one all-auction sequential result beside the integrated participation profiles. The January study refresh comprises 465 integrated cases and 31 sequential cases. Matched daily realized profits, their paired differences, model sizes, forecast inputs and fallback decisions are saved. A later infeasible sequential solve retains prior physical commands and submits zero new trades; it has no previously planned future order to submit.

11.2 Operational fallback after an infeasible auction decision

Exact planned delivery does not guarantee feasible recourse after forecast updates. Limited remaining energy, power and trading capacity may conflict with the configured terminal state or remaining delivery obligations. The physical controller can continue within SOC/power limits by accepting an imbalance even when the planning model is infeasible. A pre-noon delivery shortfall is executed and settled under the retained policy; those closed intervals are excluded from the IDA3 MILP and cannot make it infeasible.

The configurable policy `execution.infeasible_auction_policy` defaults to `error`. The full-January study uses `previous_plan`. When and only when the solver confirms infeasibility, this policy uses the last feasible deterministic snapshot for the current auction's previously unsubmitted orders and remaining requested physical commands:

$$q_{a,t}^{\text{submit}} = q_{a,t}^{\text{last feasible}}, \qquad u_t^{\text{requested}} = u_t^{\text{last feasible}}. \tag{65}$$

The orders are submitted at their actual gate and priced at actual clearing prices with ordinary fee accounting; existing commitments remain fixed. The same feedback controller executes against actual PV, respects the physical limits, and settles unavoidable deviations ex post. This is a retained plan, not a new feasible optimum under the updated forecast. No recovery optimization or constraint relaxation is introduced. Subsequent selected decisions still attempt their normal reoptimization. Missing prior plans, nondeterministic snapshots and other solver failures remain errors.

Failed decision/forecast diagnostics, reused plan sources, submitted quantities and attempt counts are saved separately from successful solves; optional IIS reports use `execution.save_infea`

sibility_diagnostics. Summary tables retain fallback-day and failed-decision counts rather than omitting those dates from the comparison. The physical trajectory can finish away from the terminal SOC target of its retained plan, just as under ordinary forecast-error feedback.

The experiment-participation command runs configured profiles on a common set of delivery days $\mathcal{D}$. The ranking metric is

$$\overline{J}(P) = \frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} J^{\text{real}}(P, d), \qquad P^{\star}_{\mathcal{D}} \in \arg\max_{P} \overline{J}(P). \tag{66}$$

It is an empirical ranking on $\mathcal{D}$, not an established optimum for unseen days. Selecting a different winning profile after each observed day would evaluate $|\mathcal{D}|^{-1} \sum_d \max_P J^{\text{real}}(P, d)$ and would use additional hindsight; the saved daily-winner table is diagnostic and is not the fixed-profile ranking.

Every profile must complete exactly the same requested dates with fully priced actual settlement. Missing or invalid inputs fail the batch; they do not silently change a profile's averaging sample. Battery and operator fee initial state are identical across profiles and reset independently each study day. This is not a continuous multi-day carry of SOC or monthly traded volume. The initialization weather cutoff is common before the earliest candidate gate and delivery start. Later forecast updates occur only at selected auction decisions. Delivery grids, initial PV forecasts, effective realized PV (including configured fallback intervals) and shared actual auction-price curves are fingerprinted and checked for equality across profiles. The market adapter resolves native auction-product durations for each date. Actual price snapshots, model inputs and all daily decisions remain in the per-profile run artifacts.

Summary outputs retain average, total, minimum, maximum and median daily profit, 10th/90th percentiles, daily population standard deviation, negative-profit and zero-active-auction day counts, daily wins including configured tolerance ties, paired differences to the largest selected profile, mean SOC/energy, curtailment, wear, transaction fees and active auction/solve counts. Final stored energy is reported without salvage valuation. The current configured terminal SOC remains a planning constraint rather than a guarantee on feedback's final state, so cash profit and ending inventory must be assessed together. Fee and inventory reset assumptions also matter when extrapolating to a longer investment horizon.

The comparison figure places a separate gold perfect-foresight reference above the ranked profiles. Its mean and daily standard deviation use the full-day perfect-information objective values (Gurobi ObjVal) on the identical dates, effective realized PV, assets, costs and configured terminal SOC, with every configured auction that remains available. This reference includes markets outside a restricted candidate/profile family. Existing all-auction benchmark outputs are reused; if that profile is absent, one additional perfect-information solve per day supplies the reference without another auction replay. Its daily objectives and summary are retained separately, along with objective-minus-profit differences in the profile tables. This is a model-objective reference rather than an additional feedback participation rank: actual imbalance settlements and realized terminal inventory can differ between the accounting cases. participation_study.perfect_foresight_reference enables this feature by default and can be set to false.

The batch YAML section participation_study controls candidate/required auctions, explicit profile lists, common dates, output paths, per-profile figures/models, summary plotting and profit-tie tolerance. The supplied example is configs/full_experiments/gme_participation_2025-01-01_03.yaml. The full-month independent-day configuration is configs/full_experiments/gme_participation_2025-01.yaml, with 31 common delivery days and 465 profile/day cases. Fully cancelled IDA3-only cases remain in the averaging sample rather than changing that profile's denominator. The driver is src/simulation/participation_study.py, while src/participation.py validates/enumerates subsets. Summary CSV/JSON/plots are stored apart from each individual replay under its profile/day directory. Optional participation_study.resume verifies the same effective configuration, completed child manifests, input fingerprints and saved

profits before reusing checkpoints. Incomplete children are retained separately and rerun; changed configurations are rejected. After independent solution validation, tolerated negative floating-point command roundoff is floored to zero before the physical controller's strict nonnegative-input checks.

12 Continuous operation with daily optimization horizons

12.1 Delivery books and the chronological event sequence

The separate `experiment-continuous` workflow selects inclusive local delivery dates $d \in \mathcal{D} = \{d_0, \dots, d_F\}$ through `continuous_optimization.start_date` and `end_date`. Each day has its own physical grid $\mathcal{T}_d$, native auction products, actual price curves and auction calendar. Commitments and revisable plans are keyed by delivery day as well as auction and interval, for example (d, a, t); the same auction name on another delivery day is a different market book. This is continuous chronological execution with daily planning horizons; it is not a single multi-day MILP, and a later-day solve cannot revise an earlier delivery book.

All events are merged into one chronological sequence. On an ordinary operating day d, the replay submits IDA3 orders for d, then DA, IDA1 and IDA2 orders for $d+1$ while executing delivery d. Submissions use the configured decision buffer before the local gate; price publications book already submitted orders without triggering another optimization. IDA3 variables remain restricted to noon–midnight of their delivery day. Native DA-hour equalities and quarter-hour IDA/physical controls are unchanged. Cancelled auctions produce no order, variables or solve, and unexplained missing inputs fail preflight.

This is a chronological sequence of *daily* MILPs, not a simultaneous month-long or two-day trading MILP. An optimization for $d+1$ cannot change current-day dispatch or commitments for d. Conversely, current-day IDA3 cannot reopen submitted next-day DA quantities. The current-day commands and book enter next-day planning through the projected boundary state below.

12.2 Forecast-based projection of the next initial state

At a decision cutoff c_k, let x_k^{obs} be the most recent observed physical state at boundary $\tau_{i_k} \le c_k$. This state includes energy, cumulative executed throughput/degradation and, when relevant, depth-model segment history. The current continuous controller requires zero configured telemetry latency, treating an interval's actual state as observable at its end; preflight rejects a nonzero `source.telemetry_latency_minutes`. This restriction is independent of the weather forecast's availability lag. An unfinished interval's already issued command is retained. Let $\bar{u}_i$ be the retained current-day charge, discharge and PV-curtailment commands and Q_i^{sub} the cumulative submitted grid position. For a pre-delivery DA, IDA1 or IDA2 solve for d, the planning boundary $b(d,a)$ is midnight at the start of d; for IDA3 it is noon on d. Define a shadow state using only the weather/PV vintage available at c_k:

$$\widehat{x}_{i_k|k} = x_k^{\mathrm{obs}}, \tag{67}$$

$$\widehat{x}_{i+1|k} = \Phi\Big(\widehat{x}_{i|k}, \bar{u}_i, Q_i^{\mathrm{sub}}/\Delta_i, \widehat{P}_{i|c_k}\Big), \qquad i_k \le i < b(d,a). \tag{68}$$

Φ is the same grid-tracking and physical-protection controller used for actual execution. Thus the forecast projection respects SOC, battery power and POI limits while preserving the old commands and submitted grid target. It uses no new optimization for the locked prefix and no future measured PV. Forecast eligibility uses the UTC publication/availability clock at c_k, including the configured six-hour lag, independently of the later projected delivery boundary.

The daily MILP fixes its initial state to $x_{d,b} = \widehat{x}_{b(d,a)|k}$. At a later gate this projection is recomputed from the newer actual state and latest eligible forecast. Meanwhile actual execution advances

separately:

$$x_{i+1}^{\text{real}} = \Phi\Big(x_i^{\text{real}}, \bar{u}_i, Q_i^{\text{sub}}/\Delta_i, P_i^{\text{real}}\Big). \tag{69}$$

Neither solving nor revising a shadow projection overwrites actual energy or aging state. The realized midnight state may differ from the previously projected state; later solves carry the actual state as it becomes available. Saved boundary records retain prediction, actual state and signed error $E_b^{\text{real}} - \widehat{E}_{b|k}$. Midnight projections are distinguished from IDA3 noon projections in diagnostics and plots.

12.3 Daily free decisions, initial condition and final target

For integrated planning, the free trade quantities are those of the current and remaining active auctions *for that delivery day*. For sequential `current_auction_only` planning, only the current auction has free trade variables and expected trading revenue. Previously submitted quantities for the target book, including unpublished pending orders, remain fixed in $Q_{d,t}^{\text{fix}}$ in both modes. Their cash values remain external to the new decision objective. Physical controls and SOC over the target day's remaining eligible delivery horizon remain free under the existing balances, costs and asset constraints. Trades in other delivery books are not decision variables in this MILP. Its objective is the daily residual objective (24); projected costs already locked before its planning boundary are not new trading opportunities.

For the supplied January configuration, $d_0 = 1$ January 2025 and $d_F = 31$ January 2025. The supplied boundary is

$$E_{d_0,0}^{\text{real}} = 0.90\, E^{\text{nominal}}. \tag{70}$$

Trading starts with the 31 December DA/IDA1/IDA2 events for delivery on 1 January. Their initial midnight state is this supplied assumption; it is not inferred from unavailable 31 December physical operation. Eligible 31 December forecasts and the actual 1 January delivery-price curves are still required. After the first midnight, there is no daily SOC reset.

The configured terminal constraint is applied only when planning the final delivery day and only when the target is non-null:

$$e_{d_F,|\mathcal{T}_{d_F}|} = E^{\text{nominal}}\, s_F/100, \qquad s_F = 50 \quad \text{(SOC percentage in the January example)}. \tag{71}$$

Intermediate midnight states are free within their SOC box; there is no intermediate target or energy salvage price. Actual terminal energy remains an output of feedback execution and need not equal the planned target. The last delivery interval ends at local midnight on 1 February. No orders for 1 February delivery are submitted by the January study.

12.4 Independent trajectories, accounting and restart

The default experiment executes three independent chronologies: integrated planning, sequential current-auction-only planning and an integrated perfect-PV reference with the same daily horizons. All start from the same supplied state and use matched actual price, settlement, effective measured PV and asset inputs, but each carries its own decisions and realized state. The reference replaces forecast PV by effective realized PV, including in shadow projections. It does not introduce future delivery-day trade variables. Its reported performance is continuous settled operating profit, not the sum of overlapping auction objective values. The two trading modes and this reference are independently configurable.

The reference defaults to one optimization per delivery day, at DA (or the first active auction if DA is cancelled). The flag `continuous_optimization.perfect_pv_reoptimize_at_each_auction` defaults to `false`; setting it to `true` restores optimization at every active gate. With the default,

all later auction quantities are read from the initial integrated plan and submitted only at their own gates, while its physical commands remain in force throughout delivery. Publication, fees, actual state carryover and settlement follow the same chronological rules. The event trace identifies each retained submission separately from solver attempts. Perfect information preserves the initial plan's attainable value; repeated solves can select different schedules when several optima exist, and finite MIP tolerance can permit small objective improvements.

Monthly operating profit is additive over delivery products:

$$J_{\mathcal{D}}^{\text{real}} = \sum_{d \in \mathcal{D}} \left[\sum_{a,t} \pi_{d,a,t}^{\text{real}} q_{d,a,t}^{\text{sub}} + \sum_{t \in \mathcal{T}_d} S_t - D_d^{\text{executed}} - C_d^{\text{import}} - C_d^{\text{trading,real}} \right]. \tag{72}$$

The trade ledger preserves submission/publication clocks; daily profit is attributed to delivery date. Fees and auction revenues reconcile both by day and over the complete ledger. Operator fee volume still uses the submission month: the first day's 31 December orders belong to December, even though their revenue appears in 1 January delivery profit. Actual fee mode therefore requires an applicable December tariff and a timestamped initial operator volume state; January volume starts according to the fee calendar, without a daily reset. Wear is charged only on actual executed throughput/history. Per-day initial expectations and realized-minus-expected differences are reported separately; no objective values across overlapping gate solves are added to obtain profit.

With the configured infeasibility policy, a later failed solve can retain the previous physical plan and its already planned current order in integrated mode; the sequential mode retains dispatch and submits zero new trades. No prior plan means the existing failure behavior applies. Actual feedback still enforces physical limits and settles remaining deviations. Failed attempts and fallback choices remain explicit diagnostics.

Resume checks the effective configuration, code/policy and common input fingerprints before restoring an atomic completed-midnight checkpoint. The checkpoint contains actual physical/aging state, outstanding priced and unpriced commitments for both relevant books, retained commands, fee state, publication status and the event cursor. Work after that boundary is replayed without double-booking cash or fees. Business decisions rolled back for resume do not survive as submitted history; every actual Gurobi attempt, including discarded work, survives in a separate attempt log with its rollback marker. Native `Model.Runtime` statistics include these attempts, keep reference solves identifiable, and report totals plus mean/minimum/maximum/population standard deviation by strategy and delivery day. Model-build and Python wall times remain separate. Resolved configurations, common-input provenance, daily/monthly KPIs, state projections and comparisons are saved by file type under `outputs/continuous_optimization/<run_id>/`.

13 Price Forecasting Status

The active GME and continuous historical experiments use realized auction clearing curves as an oracle input. They do not use predicted prices or model price uncertainty. The former retrospective regression, fundamentals, Gaussian calibration, and prepared stochastic-price input implementations are archived outside the repository in `/home/rghetti/price_prediction_archive`; they are not part of the active optimization pipeline and should not be interpreted as a validated forecasting method. The current work needed is to define price forecasts and uncertainty scenarios from information demonstrably available by each auction gate, then integrate those causal distributions into the market decisions.

14 Optional optimized-imbalance formulation

This section describes supported code that is disabled in the current farm-910 study. When `optimize_imbalances=true`, replace (23) by

$$\delta^+_{s,t} - \delta^-_{s,t} = g_{s,t}\Delta_t - Q_{s,t}, \tag{73}$$

$$0 \le \delta^+_{s,t} \le M_t y_{s,t}, \tag{74}$$

$$0 \le \delta^-_{s,t} \le M_t(1 - y_{s,t}), \qquad y_{s,t} \in \{0,1\}. \tag{75}$$

Here $M_t = (P^{\text{imp}}+P^{\text{exp}})\Delta_t$, optionally reduced by a configured imbalance-energy cap. The common M_t is valid because physical and contractual net energy use the same directional POI range. The sign binary prevents simultaneous positive/negative deviations, including with negative prices. Add to each scenario's objective

$$\sum_t \left[\lambda^+_{s,t}\delta^+_{s,t} - \lambda^-_{s,t}\delta^-_{s,t} - \kappa(\delta^+_{s,t} + \delta^-_{s,t})\right]. \tag{76}$$

If this mode has zero settlement price, zero penalty and no useful delivery restriction, an optimized contract need not be backed by delivered energy. That is why the present experiment disables optimized deviations and computes physical mismatches only afterwards. The optional formulation is a mathematical study mode, not an assertion that deliberate imbalance strategies are permitted.

15 Data provenance, outputs and reproducibility

Farm inputs are local read-only Zarr/Parquet stores. API and workbook records use `data/api/market.duckdb`. Raw/normalized snapshots, source identities, checksums, revision and retrieval metadata support repeated offline studies. Historical publication remains a separate field/convention. Provider acquisition is an explicit user-run command; neither the MILP builder nor the historical experiment downloads missing inputs. DuckDB writers and analysis processes must be coordinated; complete cached days can be studied before an entire acquisition window is complete.

YAML controls study date, calendar validity, source mappings, PV and data policies, BESS and POI characteristics, state, degradation, solver and outputs. One include is resolved relative to the containing YAML, with recursive mapping overrides. Every shipped entry point imports `configs/standard.yaml` directly; experiment YAMLs contain only their overrides. The standard exposes shared feature flags to existing experiments. Explicit null values and replacement lists survive the merge. The `resolve-config` command prints the effective YAML without running a study or accessing provider data.

A single root `farm_id`, optionally overridden with `-farm-id`, selects forecast and physical plant identifiers through typed references resolved after merging. Each farm uses its own weather and measured-power series and Parquet metadata for PV geometry and tracking. Automatic directional POI ratings use declared AC nameplate capacity divided by 1000 as an illustrative MW proxy; per-farm ratings or explicit plant limits override that rule. Battery size and SOC percentages remain independent of farm selection. The reusable `configs/full_experiments/gme_day.yaml` exposes this workflow.

The current 50% study is configured by `configs/full_experiments/gme_oracle_2025-01-01_farm910_perfect_pv_comparison_terminal_soc_50.yaml`. Its inherited assumptions include a 5 MWh battery, 2.5 MW symmetric power, a configured SOC box, 0.97 combined efficiencies and a 6.638 MW directional POI. These are reproducible study settings, not confirmed anonymous-farm hardware.

Each historical run is grouped under `outputs/full_experiments/gme_oracle/<run>/farm_<id>/<day>/` with `plots/`, `csv/`, `json/`, and optional `models/`. The effective configuration, input

provenance and artifact manifest accompany auction plans, immutable ledger, actual dispatch, used PV vintages, settlement and objective comparisons. The main KPI distinctions are:

- `total_gain_eur`: actual feedback profit after available settlement;
- `perfect_information_objective_eur`: achieved full-information PV+BESS objective;
- `pv_only_perfect_information_objective_eur`: achieved optional PV-only objective;
- `battery_incremental_operating_value_eur`: the difference in (61);
- initial expectation, realized-minus-expected revenue/profit, and separate executed versus counterfactual degradation components.

Figures preserve interval widths and both available boundary points, including next-day local midnight. Held hourly/quarter-hour prices and powers do not create new samples; SOC uses the actual boundary states. Snapshot settlement diagnostics include their different delivery windows and common-horizon comparisons.

Solver settings use relative MIP gap 10^{-4} (0.01%) by default, with no absolute MIP gap configured. Configurable `feasibility_tolerance` and `integrality_tolerance` map to Gurobi's numerical feasibility and integer tolerances. The historical YAML sets both to 10^{-8} to avoid tiny binary-mode leakage while preserving the independent validator's physical checks. Objective comparisons use achieved `ObjVal`. Status, relative gap, model sizes and timing remain recorded. Before a schedule is used, independent validation recomputes physical balances, limits, native product consistency, information sharing, degradation and objective components. Execution separately checks actual SOC and power/POI protection. Software consistency does not resolve uncertain farm metadata, forecast bias or historical publication vintages.

16 Current limitations and extension boundaries

The current empirical workflows use daily perfect-price planning with one weather member, either as independent days or consecutive continuous operation, with approximate geometry and an idealized interval-average controller. Measured AC is effective observed production, not an independently estimated unconstrained irradiance-to-generation potential; hidden plant curtailment/faults in source production are not identified automatically. Configured cleaning and fallback choices remain part of the experiment.

The model currently omits continuous trading, order acceptance, market impact, multiple-farm network coupling, ancillary services, intrainterval dynamics, ramping/startup constraints, dynamic capacity loss, calendar aging, thermal battery dynamics, fixed operating costs and investment cash-flow optimization. The generic `backtest` CLI remains synthetic-only; the separate historical `experiment-gme-oracle`, `experiment-participation` and `experiment-continuous` workflows are the implemented manual-data replays. Future uncertain-price evaluations need their own causal inputs and comparison protocol rather than relabeling the current oracle results.

The original implementation-status report is retired. This source and its PDF are the maintained mathematical description; run manifests and saved solver results are the numerical evidence for particular configurations. Practical commands and configuration examples are in the root `README.md`.

References

[1] B. Xu, J. Zhao, T. Zheng, E. Litvinov and D. S. Kirschen, *Factoring the Cycle Aging Cost of Batteries Participating in Electricity Markets*. Author preprint and published paper links retained by the repository: https://arxiv.org/abs/1707.04567 and https://doi.org/10.1109/TPWRS.2017.2733339.

[2] pvlib Python documentation. The installed library implements the solar-position, single-axis tracking, isotropic transposition, Faiman and PVWatts functions called by `forecasting/pv.py`. https://pvlib-python.readthedocs.io/.

[3] GME market timing references and period-specific study provenance are recorded in `markets.py` and the selected study's YAML. The January historical configuration explicitly labels assumed result-publication clocks.

[4] Repository source files listed in Section 1 and the effective configuration saved with each run are authoritative for implemented behavior. Conceptual references do not replace inspection of those inputs and policies.